

# ⚡ AI Cyber Threat Landscape

TLP:GREEN — INTERNAL USE

Weekly Executive Intelligence Brief

Week of 12–15 May 2026


Based on open source intelligence (OSINT) · Prepared by an AI intelligence analyst

Issued: 15 May 2026, 10:07 EDT

Prepared by: AI Intel Analyst

## ▣ RISK LEVEL REFERENCE GUIDE


## ▣ CURRENT CYBER THREAT LEVEL



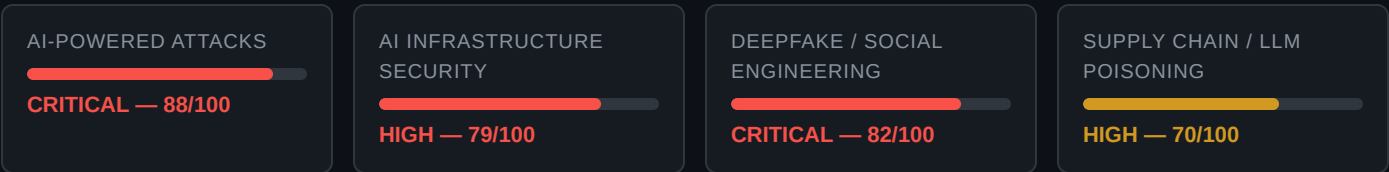
### ▲ OVERALL RISK ELEVATED

AI-powered offensive tooling has crossed a capability threshold this week. **Agentic attack frameworks operating without human-in-the-loop** are confirmed in at least two nation-state campaigns. Concurrently, the AI infrastructure layer itself is under attack — an actively exploited **Microsoft SSO Plugin zero-day** (CVSS 9.1) is compromising Jira and Confluence environments at scale, and a **Palo Alto PAN-OS zero-day** (CVE-2026-0300) linked to a China-nexus actor has been weaponized since April. The merger of AI-as-weapon and attacks-on-AI defines the current threat moment.

## ▣ THIS WEEK AT A GLANCE



## ▣ RISK POSTURE BY DOMAIN



## ▣ AI AS A WEAPON — ACTIVE THREAT VECTORS



Nation-state and criminal actors are deploying **autonomous AI agents** that execute full intrusion lifecycles — recon, exploitation, lateral movement, and exfiltration — without human operators in the loop. MuddyWater (Iran-MOIS) used an AI-assisted campaign in April 2026 masquerading as Chaos ransomware. Chinese-nexus TGR-STA-1030 is actively probing 155 countries using agentic tooling. **Time-to-exploit has compressed to hours.**

Nation-State

Active Exploitation

Autonomous

Deepfake audio/video involved in **>30% of high-impact corporate impersonation attacks** in 2025 (Cyble). Modern deepfake video **bypasses detection tools with >90% accuracy**. [SheByte PhaaS](#) platform delivers AI-generated templates at scale. Scattered Spider + Russian ransomware partnerships now combine Western social engineering with Eastern technical capability. **A \$25M Hong Kong fraud remains the benchmark.**

BEC / Fraud

\$25M+ incidents

PhaaS



### AI-Generated Phishing at Scale

Spear-phishing at spray-and-pray volume

LLMs now produce **hyper-personalized, multilingual phishing lures at mass scale**. ClickFix browser-based attacks (AI-generated fake CAPTCHAs tricking users to execute code) are escalating rapidly. AI-adaptive lures adjust tone and content mid-interaction. **Traditional awareness training is no longer sufficient** — the content is indistinguishable from legitimate communications without process controls.

Phishing

ClickFix

Browser-Based



### Ransomware 5.0 — Autonomous Operations

Fully autonomous ransomware PoCs confirmed

The next evolution: AI embedded in **every stage of the ransomware chain**. AI generates access, maps networks in minutes, chains exploits on the fly, refactors malware in real-time to evade signatures, and runs extortion negotiations autonomously. PoC frameworks are public. **TeamPCP released Shai-Hulud worm source code this week**, demonstrating AI-guided self-propagation. A semi-autonomous campaign is no longer a prediction — it is a timeline question.

Ransomware

Public PoC

Autonomous

## ATTACKS ON AI SYSTEMS — INFRASTRUCTURE UNDER SIEGE



### Prompt Injection (OWASP LLM01)

73% of production deployments remain exposed

Prompt injection is the #1 LLM vulnerability per [Zylos AI](#). **Indirect injection** — via poisoned documents, web content, and RAG pipelines — now accounts for **>55% of observed incidents**, with 20–30% higher success rates than direct attacks. **73% of production AI deployments remain exposed**. The EchoLeak attack against M365 Copilot demonstrated zero-click corporate data exfiltration.

OWASP LLM01

RAG Poisoning

Zero-Click



### AI Supply Chain & Model Poisoning

250 malicious docs = backdoor in any LLM

[Research](#) (Anthropic / UK AI Security Institute / Alan Turing Institute, Oct 2025) confirms **as few as 250 malicious documents** can backdoor models from 600M–13B parameters. OpenAI hit by TanStack supply chain attack this week. ShadowLogic backdoors survive format conversion and fine-tuning. Hugging Face model downloads remain unverified in most enterprise deployments. **MCP server misconfigurations** enable tool poisoning, credential exfiltration, and code execution at the agent layer.

Supply Chain

Backdoor

MCP Risk



### Jailbreaking & Guardrail Bypass

Policy Puppetry bypassed all major frontier models

**Policy Puppetry** — a universal jailbreak technique — successfully bypassed safety guardrails across all major frontier models in 2026. Per [Zylos AI 2025 systematic research](#): roleplay-based attacks achieve **89.6% success rate**; logic trap attacks 81.4%. iMIST (multi-step tool-disguised jailbreak) uses RL to iteratively erode safety



### CVE-2026-7482 — Bleeding Llama

Ollama heap OOB read — CVSS 9.1 — PATCHED

**Bleeding Llama** (CVE-2026-7482) — a heap out-of-bounds read in Ollama AI inference server — allows unauthenticated attackers to exfiltrate **API keys, environment variables, and chat session data** from AI inference endpoints. Affects all Ollama versions <0.17.1. Most enterprise AI lab deployments ship Ollama on default (unauthenticated)

barriers across multi-turn conversations. Guardrails are necessary but **cannot be the only control** — they are routinely bypassed.

Guardrail Bypass 89.6% ASR Multi-Turn

configuration. Patched in v0.17.1. **Prioritise patching any internal Ollama instance immediately.**

Patched CVSS 9.1 API Key Theft

CRITICAL VULNERABILITIES — WEEK OF 12–15 MAY 2026

| CVE / ID       | PRODUCT                           | TYPE                                  | CVSS      | STATUS                 | AI RELEVANCE                                            |
|----------------|-----------------------------------|---------------------------------------|-----------|------------------------|---------------------------------------------------------|
| CVE-2026-42897 | Microsoft Exchange OWA            | XSS → RCE (Zero-Day)                  | —         | UNPATCHED              | AI BEC pivot via compromised email infrastructure       |
| CVE-2026-41103 | MS SSO Plugin — Jira / Confluence | Auth Bypass / Identity Forgery        | 9.1 CRIT  | ACTIVELY EXPLOITED     | AI knowledge base + issue tracker takeover              |
| CVE-2026-20182 | Cisco Catalyst SD-WAN             | Auth Bypass — Full Admin              | 10.0 CRIT | CISA KEV — May 17 FCEB | Network segmentation bypass for AI workload isolation   |
| CVE-2026-41651 | PackageKit (Linux)                | LPE / TOCTOU — Pack2TheRoot           | 8.8 HIGH  | ACTIVELY EXPLOITED     | LPE on AI inference / training servers                  |
| CVE-2026-40402 | Windows Hyper-V                   | EoP / Guest-to-Host Escape            | 9.3 CRIT  | Patched May PT         | Cloud AI workload isolation breach                      |
| CVE-2026-0300  | Palo Alto PAN-OS                  | Unauthenticated RCE (Buffer Overflow) | CRITICAL  | Patched May 13         | Perimeter breach for AI DMZ networks; China-nexus actor |
| CVE-2026-7482  | Ollama AI Inference Server        | Heap OOB / Data Exfiltration          | 9.1 HIGH  | Patched v0.17.1        | Direct AI infrastructure attack — API key theft         |

THE AI ARMS RACE

OFFENSIVE AI CAPABILITIES

- Autonomous recon & exploitation chains
- Hyper-personalized phishing at spray-and-pray scale
- Real-time deepfake voice/video (undetectable by 2026)
- Adaptive malware that refactors against AV signatures
- AI-optimized psychological operations
- Jailbreak automation (89.6% success rates)
- Ransomware 5.0 — fully autonomous negotiation bots
- Model backdooring via <250 poisoned documents

VS

DEFENSIVE AI CAPABILITIES

- Behavioral UEBA detecting anomalous AI agent actions
- AI-powered SOC alert triage and prioritization
- NVIDIA NeMo Guardrails (1.4× detection improvement)
- Automated AI red teaming (24/7 continuous)
- Runtime inference monitoring for prompt injection
- AI-driven deepfake detection (accuracy degrading)
- Provenance-tagged RAG document pipelines
- Model signing + supply chain verification

⚠️ *Analyst Assessment:* Offensive AI is currently outpacing defensive AI in adaptability and scale. The Q3 2026 projection is analyst judgement based on current capability trends — not a sourced forecast.

📁 KEY INCIDENTS & INTELLIGENCE — THIS WEEK

📁📁 **MuddyWater — AI-Assisted Intrusion Masquerading as Ransomware** Nation-State May 2026

Iran-linked MuddyWater (MOIS) conducted an AI-assisted intrusion using Chaos ransomware as a false flag. Operations included AI-assisted recon, credential harvesting, and data theft with no actual ransomware deployed. **Assessment: ransomware artifacts were planted specifically to complicate attribution and mislead IR teams.** The use of fake Chaos samples suggests AI-aided obfuscation tooling.

📁 **OpenAI Hit by TanStack Supply Chain Attack** Supply Chain May 14, 2026

OpenAI's infrastructure was impacted by a supply chain compromise via the TanStack JavaScript library. This is emblematic of the broader trend: **AI vendors themselves are becoming high-value supply chain targets.** Any organization consuming OpenAI APIs or embedding OpenAI models should review their dependency trees and API key rotation posture.

📁 **TeamPCP Releases Shai-Hulud Worm Source Code** Worm / AI-Guided This Week

Criminal group TeamPCP publicly released source code for the Shai-Hulud AI-guided network worm. **Public availability will accelerate adoption by lower-skill actors.** The worm uses LLM-guided decision trees for propagation path selection. Expect derivative variants in the wild within 30 days.

✔️ EXECUTIVE RECOMMENDED ACTIONS

ACTION 01

**Patch Cisco SD-WAN & Exchange Zero-Day**

CVE-2026-20182 (CVSS 10.0) FCEB deadline is May 17 — two days. Apply Cisco patches immediately. CVE-2026-42897 (Exchange OWA) has no patch yet; apply OWA mitigations and restrict external OWA access now.

48-HOUR DEADLINE

ACTION 02

**Audit All Jira / Confluence SSO Deployments**

CVE-2026-41103 (CVSS 9.1) is actively exploited and allows complete identity forgery in Jira/Confluence. Patch to v1.3.3 and rotate all service account credentials immediately. Review audit logs for unauthorized access.

IMMEDIATE

ACTION 03

**Inventory & Patch All Ollama / AI Inference Servers**

Bleeding Llama (CVE-2026-7482) targets internal AI deployments directly. Audit all Ollama instances, enforce authentication, upgrade to v0.17.1+, and rotate any API keys that may have been exposed.

IMMEDIATE

ACTION 04

**Implement Wire-Transfer Verification Process**

Deepfake BEC is undetectable by technology alone. Mandate out-of-band voice verification (pre-shared codeword or callback to known number) for any wire transfer, vendor change, or executive-requested action exceeding a defined threshold.

THIS WEEK

ACTION 05

**Assess AI Supply Chain Exposure**

Review all third-party AI model sources (Hugging Face, etc.), MCP server configurations, and AI plugin permissions. Implement model signing verification and least-privilege agent credentials. TanStack supply chain event affects OpenAI consumers.

THIS MONTH

ACTION 06

**Deploy Runtime AI Security Monitoring**

Move beyond guardrails-only. Implement runtime inference monitoring to detect prompt injection in production AI systems. Guardrails are bypassed in 73% of tested deployments — runtime visibility is the required next layer.

Q2 PRIORITY

□ SOURCES & CITATION NOTES

| # | STAT / CLAIM                                            | SOURCE                                                                  | LINK                                                            |
|---|---------------------------------------------------------|-------------------------------------------------------------------------|-----------------------------------------------------------------|
| 1 | 73% of prod AI deployments exposed to prompt injection  | Zylos AI — LLM Security & Safety 2026                                   | <a href="#">zylos.ai</a>                                        |
| 2 | 89.6% / 81.4% / 76.2% jailbreak attack success rates    | Zylos AI — 2025 systematic research                                     | <a href="#">zylos.ai</a>                                        |
| 3 | 250 malicious documents can backdoor any LLM (600M–13B) | Anthropic / UK AI Security Institute / Alan Turing Institute (Oct 2025) | <a href="#">anthropic.com</a>                                   |
| 4 | >30% of corporate impersonation attacks used deepfakes  | Cyble Executive Threat Monitoring Report                                | <a href="#">cyble.com</a>                                       |
| 5 | >90% deepfake bypass rate against detection tools       | Cyble Deepfake-as-a-Service 2025 Report                                 | <a href="#">cyble.com</a>                                       |
| 6 | \$25M Hong Kong deepfake fraud                          | Cyble 2026 Outlook / widely reported                                    | <a href="#">cyble.com</a>                                       |
| 7 | SheByte PhaaS ~\$200/month                              | SecurityWeek — Cyber Insights 2026: Social Engineering                  | <a href="#">securityweek.com</a>                                |
| 8 | CVE data, patch status, CISA KEV entries                | CISA KEV / Microsoft MSRC / BleepingComputer May 2026 PT                | <a href="#">cisa.gov</a> · <a href="#">bleepingcomputer.com</a> |

⚠ Analyst Assessments — Not Externally Sourced

- **Risk domain scores** (88, 79, 82, 70 /100) — analyst synthesis of aggregated OSINT; not derived from a single scored framework.
- **Q3 2026 capability gap projection** — analyst trend judgement; not a sourced forecast.
- **Week-on-week prompt injection delta** — the 73% figure is sourced; the comparative baseline is an interpolated estimate.